

A Brief Introduction to R

Viswanathan Satheesh

2022-07-31

Basic operations

Some very basic operations you can carry out in R.

```
1 + 1 # addition
```

```
## [1] 2
```

```
2 - 1 # subtraction
```

```
## [1] 1
```

```
2 * 2 # multiplication
```

```
## [1] 4
```

```
6 / 2 # division
```

```
## [1] 3
```

```
3 ** 2 # exponential
```

```
## [1] 9
```

```
3 ^ 2 # exponential
```

```
## [1] 9
```

Here R works like a calculator. Apart from numbers, R can also help us print letters or a string of letters.

```
"a"
```

```
## [1] "a"
```

```
"language"
```

```
## [1] "language"
```

```
"R is my favourite programming language"
```

```
## [1] "R is my favourite programming language"
```

When working with large numbers such as **1934929292** and **23992343**, we cannot keep them in mind, or for that matter, remember complex computations. So, we have the concept of object or variable.

```
a <- 1934929292  
b <- 23992343
```

Here, we assign “<-” the first number to “a” and the second to “b”. The “<-” is called the “assignment operator”. “a” and “b” are called objects or variables. This now enables us to actually use the variables for doing further operations as seen below.

```
a + b
```

```
## [1] 1958921635
```

```
a - b
```

```
## [1] 1910936949
```

```
a * b
```

```
## [1] 4.642349e+16
```

```
a ^ b
```

```
## [1] Inf
```

We can do something similar with strings too.

```
x <- "language"  
y <- "R is my favourite programming language"
```

```
print(x)
```

```
## [1] "language"
```

```
print(y)
```

```
## [1] "R is my favourite programming language"
```

Creating a vector

In R, vectors are the most basic data objects. Let us create vectors **x** and **y**. We will do that in two ways. One, using the *c()* function, and the other using the *seq()* function. ‘c’ combines values into a vector. ‘seq’ is a sequence generator.

```
x <- c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
y <- seq(11, 20)
x
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
y
```

```
## [1] 11 12 13 14 15 16 17 18 19 20
```

The code above creates two vectors, **x** and **y**. Adding the two vectors gives:

```
x + y
```

```
## [1] 12 14 16 18 20 22 24 26 28 30
```

The elements are added element-wise. The operations in R are element-wise. As an exercise you can try doing the other mathematical operations on the two vectors.

Indexing

The elements in the vectors are indexed. So, to extract an element you need only know its position. To extract the first element in **x** and in **y**:

```
x[1]
```

```
## [1] 1
```

This returns **1**. Try the following and see what you get.

```
y[1:4]
```

```
## [1] 11 12 13 14
```

```
y[c(1, 3, 5)]
```

```
## [1] 11 13 15
```

```
y[c(-1, -3, -5)]
```

```
## [1] 12 14 16 17 18 19 20
```

```
y[-c(1, 3, 5)]
```

```
## [1] 12 14 16 17 18 19 20
```

Data types

What are the important data types? They can be listed as:

- **integer**
- **numeric**
- **logical**
- **character** Let us first create some vectors

```
n <- 1 # numeric
i <- 1L # integer
l <- TRUE # logical
c <- "Some string" #character
```

Here we have four vectors created. `x` is a numeric vector, `y` an integer vector, `t` a logical vector, and `c` is a character vector. Remember, in R everything is a vector. There are no scalars. Therefore, all these vectors that we have created are all vectors of length one. To check the length of the vector, use the `length()` function:

```
length(n)
```

```
## [1] 1
```

```
length(i)
```

```
## [1] 1
```

```
length(l)
```

```
## [1] 1
```

```
length(c)
```

```
## [1] 1
```

You can see that all objects created are of length one.

Now let us check these vectors using the `class()` function:

```
class(x)
```

```
## [1] "numeric"
```

```
class(y)
```

```
## [1] "integer"
```

```
class(t)
```

```
## [1] "function"
```

```
class(c)
```

```
## [1] "character"
```

Now let us create vectors of length > 1

```
num_vr <- c(1, 3.0, 5.0) # numeric vector  
int_vr <- c(1L, 3L, 5L) # integer vector  
log_vr <- c(TRUE, FALSE, TRUE) # logical vector  
char_vr <- c("I am", "a", "string.") # character vector
```

Now get their class.

```
class(num_vr); length(num_vr)
```

```
## [1] "numeric"
```

```
## [1] 3
```

```
class(int_vr); length(int_vr)
```

```
## [1] "integer"
```

```
## [1] 3
```

```
class(log_vr); length(log_vr)
```

```
## [1] "logical"
```

```
## [1] 3
```

```
class(char_vr); length(char_vr)
```

```
## [1] "character"
```

```
## [1] 3
```

We learnt to make vectors before, and now we have learnt to understand them a bit more. We now move on to matrices. First, let us create some vectors.

```
v1 <- 1:5  
v2 <- 6:10  
v3 <- 11:15
```

We have three vectors v1, v2, and v3 and we are going to bind them column-wise.

```
cbind(v1, v2, v3)
```

```
##      v1 v2 v3
## [1,]  1  6 11
## [2,]  2  7 12
## [3,]  3  8 13
## [4,]  4  9 14
## [5,]  5 10 15
```

The output just spews out to the console, which is not helpful. Let us create a variable, `my_mat`, and store the output

```
my_mat <- cbind(v1, v2, v3)
my_mat
```

```
##      v1 v2 v3
## [1,]  1  6 11
## [2,]  2  7 12
## [3,]  3  8 13
## [4,]  4  9 14
## [5,]  5 10 15
```

Here, we used a function, `cbind()`, to bind three vectors into three columns. Now let us use the class function on the `my_mat` variable.

```
class(my_mat)
```

```
## [1] "matrix" "array"
```

`my_matrix` is a matrix. It has three columns, `v1`, `v2`, and `v3`. And, as it should be clear now, we used three vectors to create a matrix. Let us now see an alternate method for creating a matrix.

```
trial_mat <- matrix(1:20, nrow=4, ncol=5)
trial_mat
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1    5    9   13   17
## [2,]    2    6   10   14   18
## [3,]    3    7   11   15   19
## [4,]    4    8   12   16   20
```

This creates a matrix with 4 rows and 5 columns. The `[1,]` refers to the first row. The `[,1]` refers to the first column.

Let us now talk about another kind of data structure, data frame. So, a data frame is similar to a matrix, but it can hold vectors of different classes. Let us re-create the same vectors we created in the previous session.

```

num_vr <- c(1, 3.0, 5.0) # numeric vector
int_vr <- c(1L, 3L, 5L) # integer vector
log_vr <- c(TRUE, FALSE, TRUE) # logical vector
char_vr <- c("I am", "a", "string.") # character vector

# Let us use the cbind() function to put them together.

new_mat <- cbind(num_vr, int_vr, log_vr, char_vr)
new_mat

##      num_vr int_vr log_vr char_vr
## [1,] "1"      "1"      "TRUE"  "I am"
## [2,] "3"      "3"      "FALSE" "a"
## [3,] "5"      "5"      "TRUE"  "string."

```

Looking at the output, we know that it is something we do not want. What is the class of the new variable?

```
class(new_mat)
```

```
## [1] "matrix" "array"
```

The class of the object `new_mat` is “matrix”. A matrix can hold data belonging to a particular class. In this case, every data point is converted into a character. This is called coercion. Here we need a different kind of data structure that can hold different classes of data. To demonstrate this point, let us create some vectors that we will make use of in creating this structure.

```

set.seed(1234) # since the numbers are random, this will make sure we always
               # get the same set of random numbers
plant_height <- rnorm(100, 110, 10)
head(plant_height)

```

```
## [1] 97.92934 112.77429 120.84441 86.54302 114.29125 115.06056
```

```

# Too many decimals. Let us round it off to two.
plant_height <- round(plant_height, 2)
head(plant_height)

```

```
## [1] 97.93 112.77 120.84 86.54 114.29 115.06
```

```

set.seed(237)
flowering_50 <- round(rnorm(100, 100, 10))
head(flowering_50)

```

```
## [1] 102 100 101 100 107 101
```

```

set.seed(6438)
spikelet_fertility <- round(rnorm(100, 90, 3), 2)
head(spikelet_fertility)

```

```
## [1] 89.05 89.18 93.49 92.28 88.04 93.61
```

```
max(spikelet_fertility)
```

```
## [1] 95.27
```

```
set.seed(345)
thousand_seed_weight <- round(rnorm(100, 22, 3), 2)
head(thousand_seed_weight)
```

```
## [1] 19.65 21.16 21.52 21.13 21.80 20.10
```

Now let us combine the four vectors into a single data structure.

```
my_data <- cbind(plant_height, flowering_50, spikelet_fertility, thousand_seed_weight)
head(my_data)
```

```
##      plant_height flowering_50 spikelet_fertility thousand_seed_weight
## [1,]          97.93          102             89.05             19.65
## [2,]          112.77          100             89.18             21.16
## [3,]          120.84          101             93.49             21.52
## [4,]           86.54          100             92.28             21.13
## [5,]          114.29          107             88.04             21.80
## [6,]          115.06          101             93.61             20.10
```

```
class(my_data)
```

```
## [1] "matrix" "array"
```

Let us now create some numbers that we will use as genotype ids. We have 100 observations and that makes it 100 genotypes. We will name the genotypes from “001” to “100”. Let us use the paste() function to create these ids.

```
a1 <- paste("00", 1:9, sep = "")
a1
```

```
## [1] "001" "002" "003" "004" "005" "006" "007" "008" "009"
```

```
a2 <- paste("0", 10:99, sep = "")
a2
```

```
## [1] "010" "011" "012" "013" "014" "015" "016" "017" "018" "019" "020" "021"
## [13] "022" "023" "024" "025" "026" "027" "028" "029" "030" "031" "032" "033"
## [25] "034" "035" "036" "037" "038" "039" "040" "041" "042" "043" "044" "045"
## [37] "046" "047" "048" "049" "050" "051" "052" "053" "054" "055" "056" "057"
## [49] "058" "059" "060" "061" "062" "063" "064" "065" "066" "067" "068" "069"
## [61] "070" "071" "072" "073" "074" "075" "076" "077" "078" "079" "080" "081"
## [73] "082" "083" "084" "085" "086" "087" "088" "089" "090" "091" "092" "093"
## [85] "094" "095" "096" "097" "098" "099"
```



```
genotypes <- c(a1, a2, 100)
genotypes
```

```
## [1] "001" "002" "003" "004" "005" "006" "007" "008" "009" "010" "011" "012"
## [13] "013" "014" "015" "016" "017" "018" "019" "020" "021" "022" "023" "024"
## [25] "025" "026" "027" "028" "029" "030" "031" "032" "033" "034" "035" "036"
## [37] "037" "038" "039" "040" "041" "042" "043" "044" "045" "046" "047" "048"
## [49] "049" "050" "051" "052" "053" "054" "055" "056" "057" "058" "059" "060"
## [61] "061" "062" "063" "064" "065" "066" "067" "068" "069" "070" "071" "072"
## [73] "073" "074" "075" "076" "077" "078" "079" "080" "081" "082" "083" "084"
## [85] "085" "086" "087" "088" "089" "090" "091" "092" "093" "094" "095" "096"
## [97] "097" "098" "099" "100"
```

Let us add this vector to our my_data object.

```
newdat <- cbind(genotypes, my_data)
head(newdat)
```

```
##      genotypes plant_height flowering_50 spikelet_fertility
## [1,] "001"      "97.93"      "102"      "89.05"
## [2,] "002"      "112.77"     "100"      "89.18"
## [3,] "003"      "120.84"     "101"      "93.49"
## [4,] "004"      "86.54"      "100"      "92.28"
## [5,] "005"      "114.29"     "107"      "88.04"
## [6,] "006"      "115.06"     "101"      "93.61"
##      thousand_seed_weight
## [1,] "19.65"
## [2,] "21.16"
## [3,] "21.52"
## [4,] "21.13"
## [5,] "21.8"
## [6,] "20.1"
```

We have seen this problem before; the entire data getting converted into a character class. To overcome this problem we use the data.frame() function.

```
field_data <- data.frame(genotypes, my_data)
head(field_data)
```

```
##      genotypes plant_height flowering_50 spikelet_fertility thousand_seed_weight
## 1      001      97.93      102      89.05      19.65
## 2      002     112.77      100      89.18      21.16
## 3      003     120.84      101      93.49      21.52
## 4      004      86.54      100      92.28      21.13
## 5      005     114.29      107      88.04      21.80
## 6      006     115.06      101      93.61      20.10
```

This output is more like it. Let us check the class of the df object.

```
class(field_data)
```

```
## [1] "data.frame"
```

It is a dataframe. A dataframe, unlike a matrix, can hold vectors of different classes. Using the most important function in R, `str()`, we get a glimpse of what the field object contains.

```
str(field_data)
```

```
## 'data.frame': 100 obs. of 5 variables:
## $ genotypes      : chr  "001" "002" "003" "004" ...
## $ plant_height   : num  97.9 112.8 120.8 86.5 114.3 ...
## $ flowering_50    : num  102 100 101 100 107 101 92 95 93 78 ...
## $ spikelet_fertility : num  89 89.2 93.5 92.3 88 ...
## $ thousand_seed_weight: num  19.6 21.2 21.5 21.1 21.8 ...
```

The “field” object is a data frame with 100 observations and 4 variables. Except for `a`, which is a factor, `plant_height`, `flowering_50`, and `spikelet_fertility` are numeric. Remember that ‘a’ is a vector containing the genotype ids. Therefore, “a” is recognised as a factor here.

Writing data to file

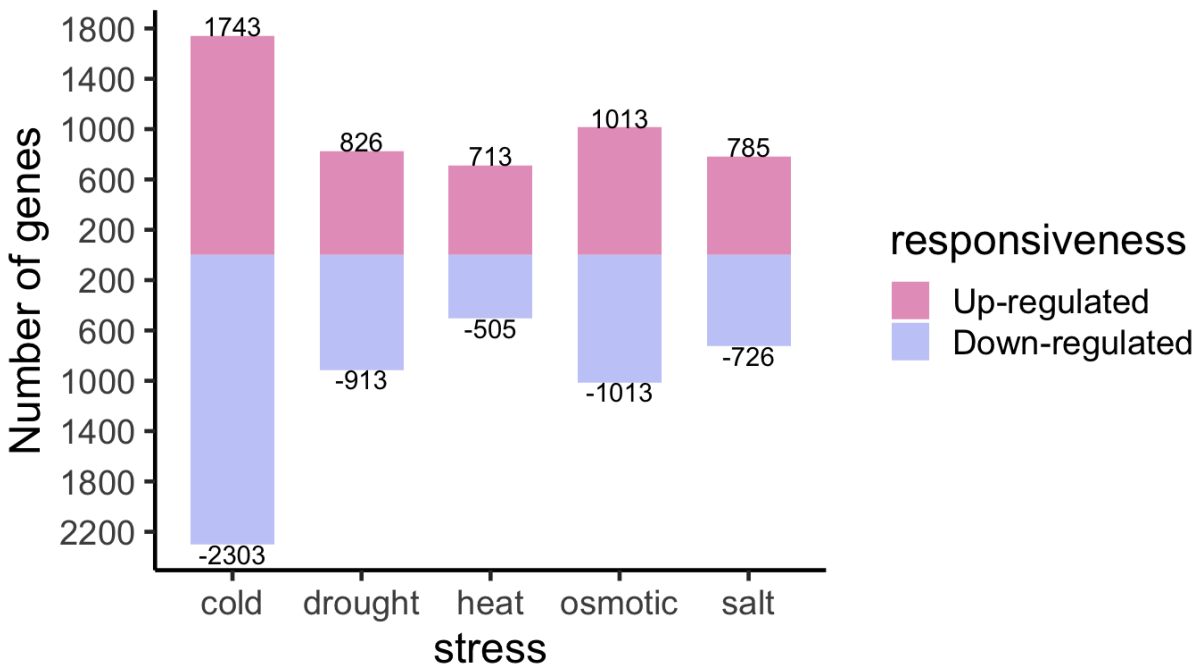
```
write.csv(field_data, "field_data.csv", quote = F, row.names = F)
```

Plotting

Creating a mirror plot

In this section we will see how to make the plot shown below.

Differential Expression of genes



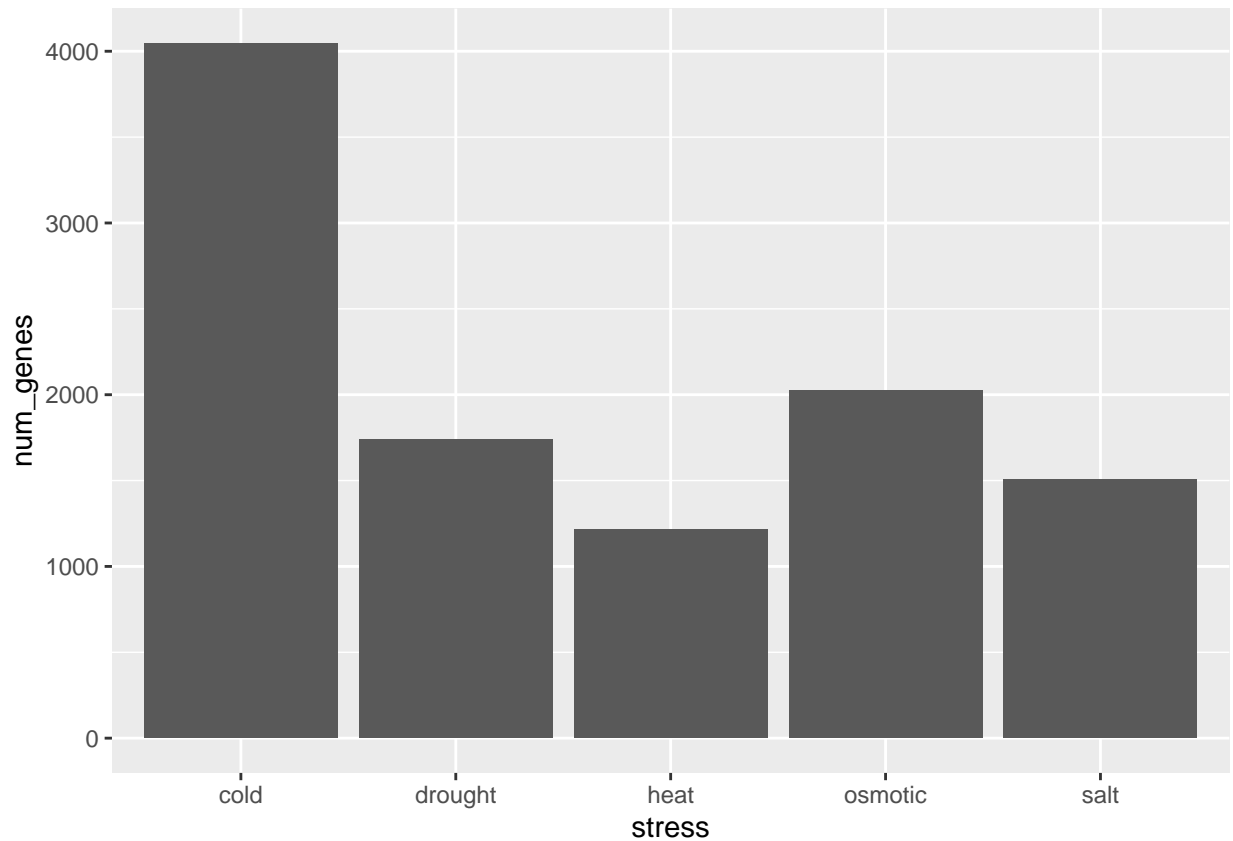
Reading data in:

```
gene_nums_mirror <- read.csv( "up_down_gene_numbers.csv" )
gene_nums_mirror
```

```
##      stress responsiveness num_genes
## 1    cold              up      1743
## 2    cold             down     2303
## 3  osmotic              up      1013
## 4  osmotic             down     1013
## 5    heat              up       713
## 6    heat             down      505
## 7    salt              up       785
## 8    salt             down      726
## 9  drought              up       826
## 10 drought            down      913
```

This data set has three columns and 10 rows. It is about differentially expressed genes under different stress conditions. So, let us start plotting with *ggplot2*.

```
library( ggplot2 )
ggplot(data = gene_nums_mirror, aes(x = stress, y = num_genes)) +
  geom_bar(stat = "identity")
```



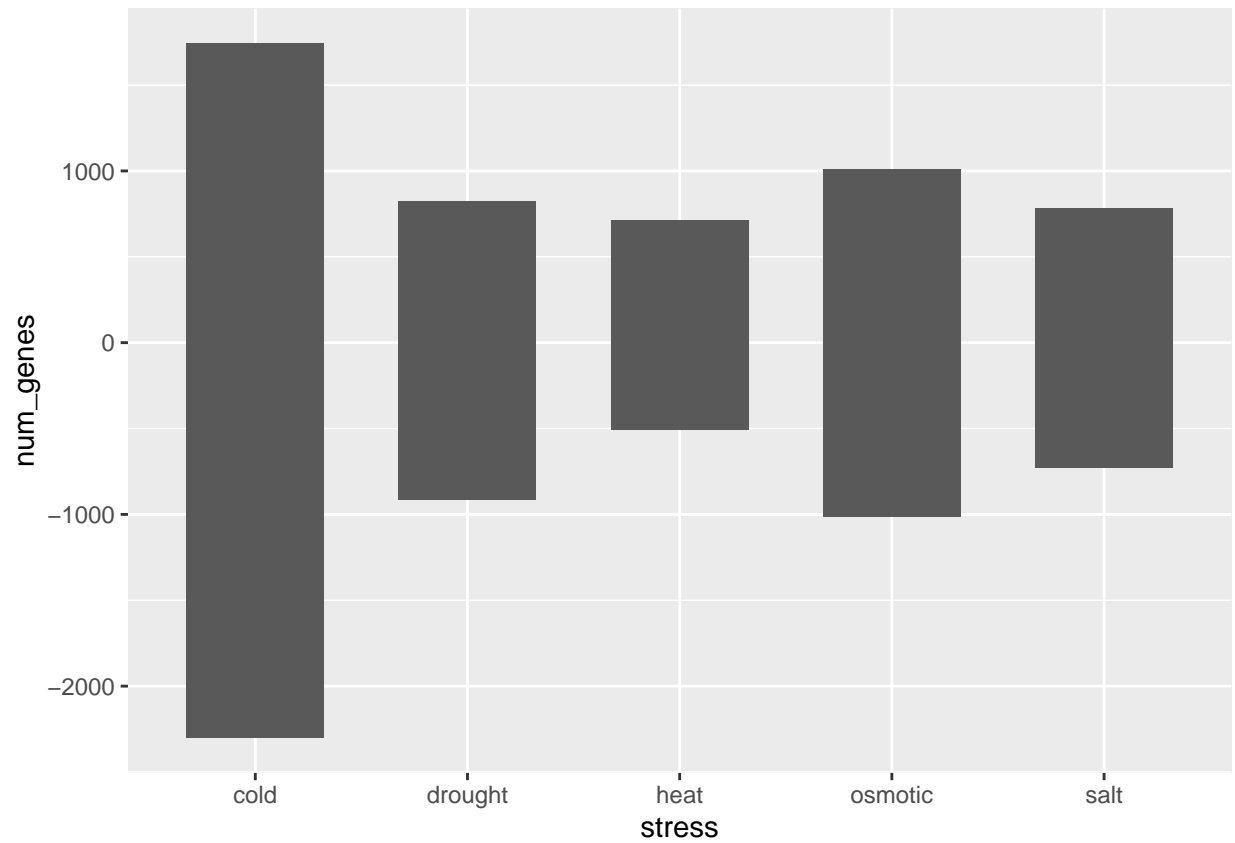
```
geom_bar( stat = "identity", position = "identity", width = 0.65 )
```

```
## geom_bar: width = 0.65, na.rm = FALSE, orientation = NA
## stat_identity: na.rm = FALSE
## position_identity
```

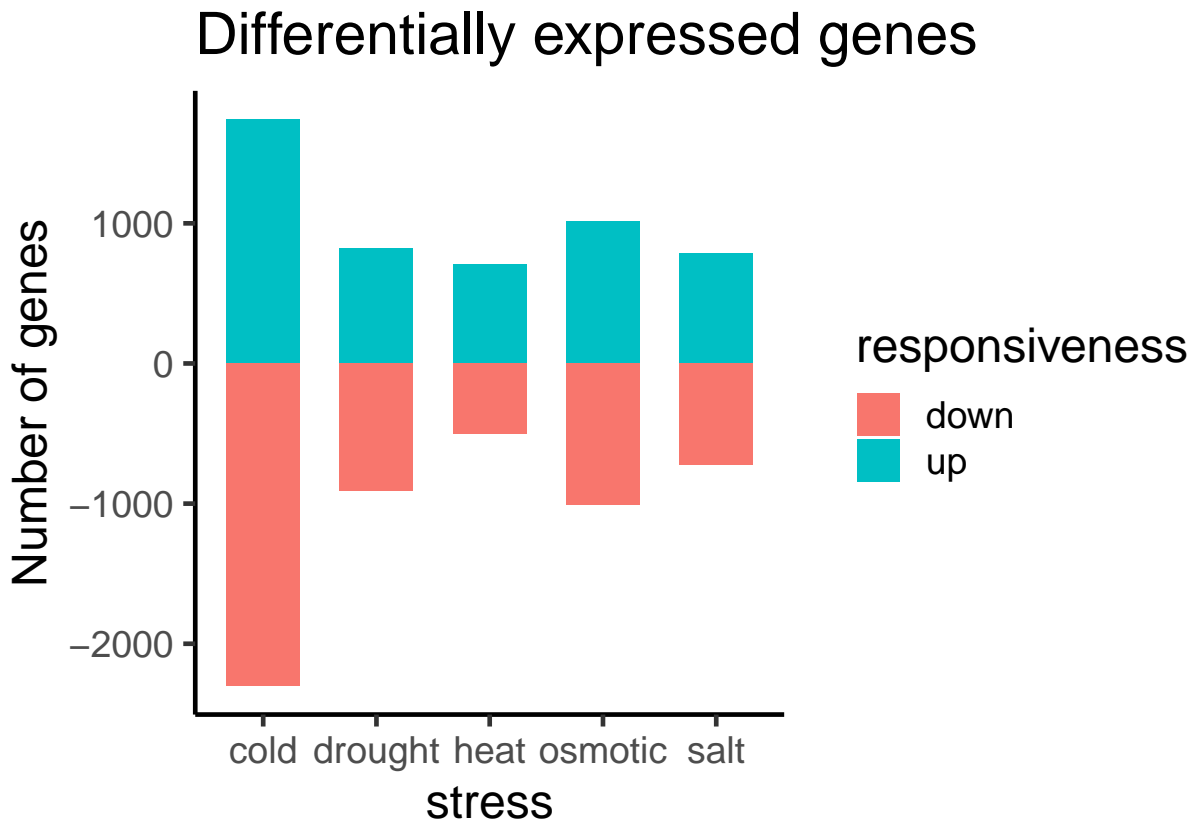
```
gene_nums_mirror$num_genes[gene_nums_mirror$responsiveness=="down"] <- -gene_nums_mirror$num_genes[gene_nums_mirror$responsiveness=="up"]
gene_nums_mirror
```

```
##      stress responsiveness num_genes
## 1    cold             up      1743
## 2    cold            down     -2303
## 3  osmotic             up       1013
## 4  osmotic            down     -1013
## 5    heat             up        713
## 6    heat            down     -505
## 7    salt             up        785
## 8    salt            down     -726
## 9  drought             up        826
## 10 drought           down     -913
```

```
ggplot(data = gene_nums_mirror, aes(x = stress, y = num_genes)) +
  geom_bar(stat = "identity", width = 0.65)
```



```
ggplot(data = gene_nums_mirror, aes(x = stress, y = num_genes,  
                                     fill = responsiveness)) +  
  geom_bar(stat = "identity", width = 0.65) +  
  theme_classic(base_size = 18) +  
  ylab("Number of genes") +  
  ggtitle("Differentially expressed genes")
```



The final code snippet:

```
library( ggplot2 )
library( wesanderson )

gene_nums_mirror <- dplyr::mutate( gene_nums_mirror,
                                   responsiveness = forcats::fct_relevel(responsiveness, "up", "down"))

ggplot(data = gene_nums_mirror, aes(x = stress, y = num_genes,
                                   fill = factor(responsiveness,
                                                  labels = c("Up-regulated", "Down-regulated")))) +
  labs(fill = "responsiveness") +
  geom_bar( stat = "identity", position = "identity", width = 0.65 ) +
  ylab("Number of genes") +
  theme_classic(base_size = 18) +
  ggtitle("Differential Expression of genes") +
  geom_text(aes(label = num_genes), vjust = ifelse(gene_nums_mirror$num_genes>0, 0,1), colour = "black",
            size = 12) +
  scale_y_continuous(breaks=seq(-3000,1800,by=400),labels=abs(seq(-3000,1800,by=400))) +
  scale_fill_manual(values = wes_palette(n=2, name = "GrandBudapest2"))
```

Differential Expression of genes

